

Section 12 Lesson 1: Working with Sequences

Try It / Solve It

1. Using CREATE TABLE AS subquery syntax, create a seq_d_songs table of all the columns in the DJs on Demand database table d_songs. Use the SELECT * in the subquery to make sure that you have copied all of the columns.

Rows

Save Run

```
create table seq_d_songs
as (select * from d_songs);
```

Results Explain Describe Saved SQL History

Table created.

0.47 seconds

Workspace: US_1350 User: US_1350_SQL01_S25

2. Because you are using copies of the original tables, the only constraints that were carried over were the NOT NULL constraints. Create a sequence to be used with the primary-key column of the seq_d_songs table. To avoid assigning primary-key numbers to these tables that already exist, the sequence should start at 100 and have a maximum value of 1000. Have your sequence increment by 2 and have NOCACHE and NOCYCLE. Name the sequence seq_d_songs_seq.

Rows

Save Run

```
create sequence seq_d_songs_seq
increment by 2
start with 100
maxvalue 1000
nocache
nocycle;
```


Results Explain Describe Saved SQL History

Sequence created.

0.00 seconds

Workspace: US_1350 User: US_1350_SQL01_S25

3. Query the USER_SEQUENCES data dictionary to verify the seq_d_songs_seq SEQUENCE settings.

Rows 

Save Run

```
select *  
from user_sequences  
where sequence_name =  
upper('seq_d_songs_seq');
```

Results Explain Describe Saved SQL History

SEQUENCE_NAME	MIN_VALUE	MAX_VALUE	INCREMENT_BY	CYCLE_FLAG	ORDER_FLAG	CACHE_SIZE	LAST_NUMBER
SEQ_D_SONGS_SEQ	1	1000	2	N	N	0	100



1 rows returned in 0.00 seconds [Download](#)

Application Express 4.2.5.00.08

Workspace: US_1350 User: US_1350_SQL01_S25 Language: en | Copyright © 1999, 2014, Oracle. All rights reserved.

4. Insert two rows into the seq_d_songs table. Be sure to use the sequence that you created for the ID column. Add the two songs shown in the graphic.

ID	TITLE	DURATION	ARTIST	TYPE_CODE
	Island Fever	5 min	Hawaiian Islanders	12
	Castle of Dreams	4 min	The Wanderers	77

Rows   Save Run



```
insert into seq_d_songs
(id, title, duration,
artist, type_code)
values (seq_d_songs_seq.nextval,
'island Fever', '5 min',
'Hawaiian Islanders', 12);
```

Results Explain Describe Saved SQL History

1 row(s) inserted.

0.02 seconds

Workspace: US_1350 User: US_1350_SQL01_S25

Rows   Save Run

```
insert into seq_d_songs
(id, title, duration,
artist, type_code)
values (seq_d_songs_seq.nextval,
'Castle of Dreams', '4 min',
'The Wanderers', 77);
```

Results Explain Describe Saved SQL History

1 row(s) inserted.

0.01 seconds

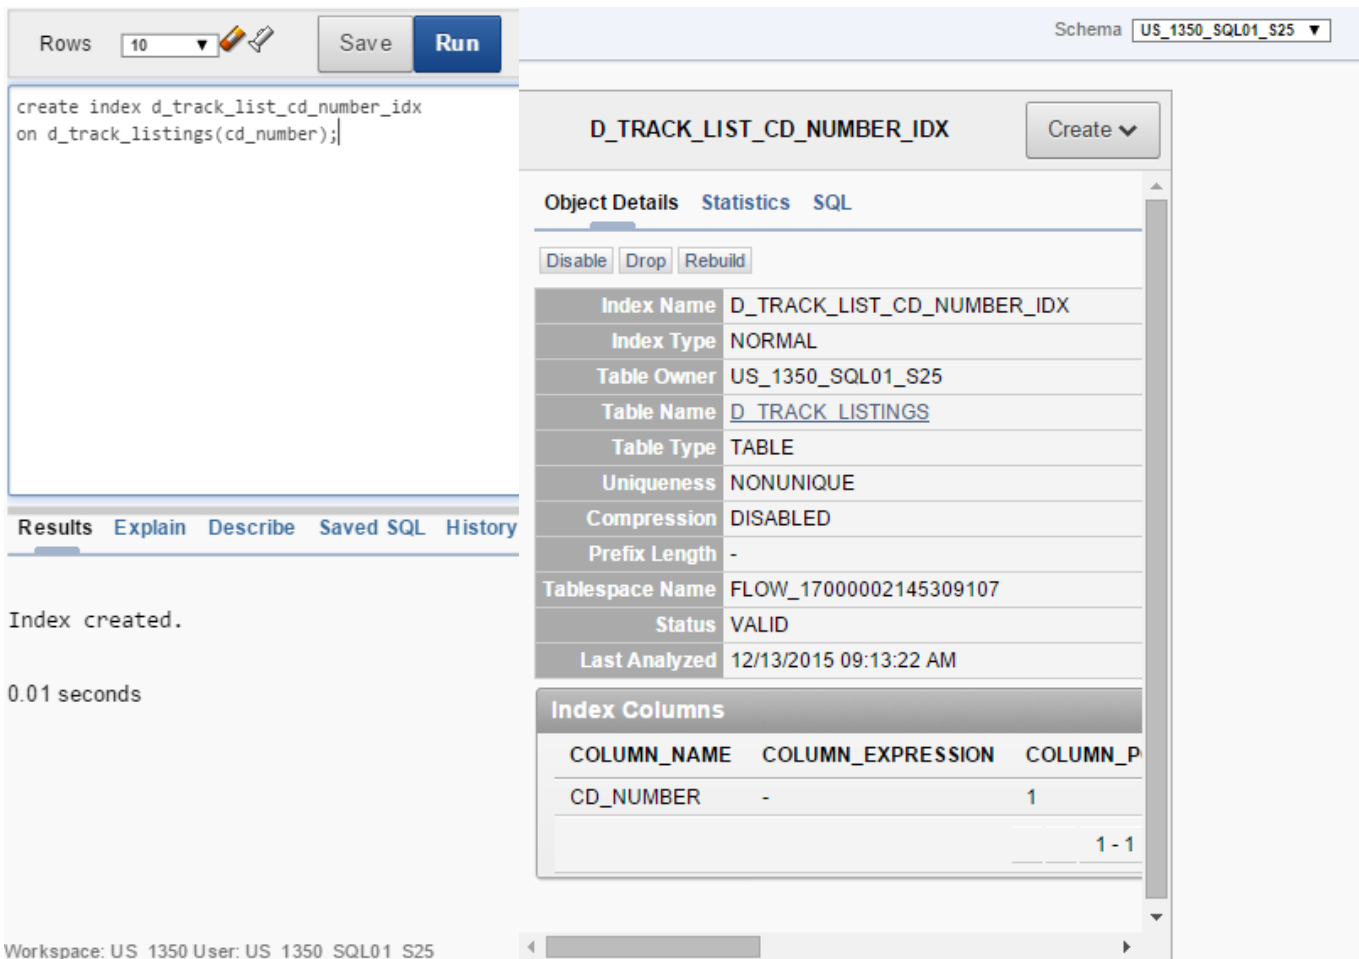
Workspace: US_1350 User: US_1350_SQL01_S25

5. Write out the syntax for seq_d_songs_seq to view the current value for the sequence. Use the DUAL table. (Oracle Application Developer will not run this query.)
- ```
SELECT seq_d_songs_seq.CURRVAL
FROM dual;
```
6. What are three benefits of using SEQUENCES?
- Sequences are time-saving, sharable, and eliminate the worry of duplicating values.
7. What are the advantages of caching sequence values?
- Cache sequences in memory provide faster access to sequence values and are populated the first time the sequence is referred to.
8. Name three reasons why gaps may occur in a sequence?
- Rolling back a statement containing a sequence, a system crash, or the same sequence being used on multiple tables.

## Section 12 Lesson 2: Indexes and Synonyms

### Try It / Solve It

1. What is an index and what is it used for? An index is a schema object that can speed up retrieval of rows by using a pointer. It can be used to provide direct and fast access to rows in a table.
2. What is a ROWID, and how is it used? A ROWID is a base 64 string representation of the row address containing block identifier, row location in the block, and the database file identifier.
3. When will an index be created automatically? An index will be created automatically when a column in a table is defined to have a PRIMARY KEY or UNIQUE KEY
4. Create a nonunique index (foreign key) for the DJs on Demand column (cd\_number) in the D\_TRACK\_LISTINGS table. Use the Oracle Application Developer SQL Workshop Data Browser to confirm that the index was created.



The screenshot shows the Oracle SQL Workshop interface. On the left, the SQL editor contains the command: `create index d_track_list_cd_number_idx on d_track_listings(cd_number);`. Below the editor, the 'Results' tab shows the message 'Index created.' and the execution time '0.01 seconds'. On the right, the 'Object Details' tab for the index 'D\_TRACK\_LIST\_CD\_NUMBER\_IDX' is displayed. The index is a NORMAL, NONUNIQUE index on the table 'D\_TRACK\_LISTINGS' in the schema 'US\_1350\_SQL01\_S25'. The index is currently VALID and was last analyzed on 12/13/2015 at 09:13:22 AM. The 'Index Columns' section shows that the index is on the 'CD\_NUMBER' column, with a column expression of '-' and a column position of 1.


| Object Name                | Object Type |
|----------------------------|-------------|
| D_TRACK_LIST_CD_NUMBER_IDX | INDEX       |

| Property        | Value                      |
|-----------------|----------------------------|
| Index Name      | D_TRACK_LIST_CD_NUMBER_IDX |
| Index Type      | NORMAL                     |
| Table Owner     | US_1350_SQL01_S25          |
| Table Name      | D_TRACK_LISTINGS           |
| Table Type      | TABLE                      |
| Uniqueness      | NONUNIQUE                  |
| Compression     | DISABLED                   |
| Prefix Length   | -                          |
| Tablespace Name | FLOW_17000002145309107     |
| Status          | VALID                      |
| Last Analyzed   | 12/13/2015 09:13:22 AM     |

| COLUMN_NAME | COLUMN_EXPRESSION | COLUMN_POS |
|-------------|-------------------|------------|
| CD_NUMBER   | -                 | 1          |

Workspace: US\_1350 User: US\_1350\_SQL01\_S25

5. Use the join statement to display the indexes and uniqueness that exist in the data dictionary for the DJs on Demand D\_SONGS table.

Rows  

Save Run

```
select ix.index_name, ic.column_name,
ic.column_position col_pos,
ix.uniqueness
from user_indexes ix
join user_ind_columns ic
on ic.index_name = ix.index_name
and ic.table_name = upper('d_songs');
```

Results Explain Describe Saved SQL History

| INDEX_NAME  | COLUMN_NAME | COL_POS | UNIQUENESS |
|-------------|-------------|---------|------------|
| D_SNG_ID_PK | ID          | 1       | UNIQUE     |

1 rows returned in 3.22 seconds [Download](#)

Workspace: US\_1350 User: US\_1350\_SQL01\_S25

6. Use a SELECT statement to display the index\_name, table\_name, and uniqueness from the data dictionary USER\_INDEXES for the DJs on Demand D\_EVENTS table.

Rows

Save Run

```
select index_name, table_name, uniqueness
from user_indexes
where table_name = upper('d_events');
```

Results Explain Describe Saved SQL History

| INDEX_NAME  | TABLE_NAME | UNIQUENESS |
|-------------|------------|------------|
| D_EVE_ID_PK | D_EVENTS   | UNIQUE     |

1 rows returned in 2.88 seconds [Download](#)

Workspace: US\_1350 User: US\_1350\_SQL01\_S25

7. Write a query to create a synonym called dj\_tracks for the DJs on Demand d\_track\_listings table.

Rows

Save Run

```
create synonym dj_tracks
for d_track_listings;
```

Results Explain Describe Saved SQL History



Synonym created.

0.01 seconds

Workspace: US\_1350 User: US\_1350\_SQL01\_S25



8. Create a function-based index for the last\_name column in DJs on Demand D\_PARTNERS table that makes it possible not to have to capitalize the table name for searches. Write a SELECT statement that would use this index.

Rows    Save Run

```
create index lower_last_name_idx
on d_partners (lower(last_name));
```

**Results** Explain Describe Saved SQL History

Index created.

0.01 seconds

Workspace: US\_1350 User: US\_1350\_SQL01\_S25

Rows    Save Run

```
select id, first_name, last_name
from d_partners
where lower(last_name) = 'cho';
```


**Results** Explain Describe Saved SQL History

| ID | FIRST_NAME | LAST_NAME |
|----|------------|-----------|
| 11 | Jennifer   | cho       |

1 rows returned in 0.01 seconds [Download](#)

Workspace: US\_1350 User: US\_1350\_SQL01\_S25

9. Create a synonym for the D\_TRACK\_LISTINGS table. Confirm that it has been created by querying the data dictionary.

Rows  

Save Run

```
create synonym d_track_list
for d_track_listings;
```

Results Explain Describe Saved SQL History

Synonym created.

0.00 seconds

Workspace: US\_1350 User: US\_1350\_SQL01\_S25

Rows  

Save Run

```
select * from
user_synonyms
where synonym_name = upper('d_track_list');
```

Results Explain Describe Saved SQL History

| SYNONYM_NAME | TABLE_OWNER       | TABLE_NAME       | DB_LINK |
|--------------|-------------------|------------------|---------|
| D_TRACK_LIST | US_1350_SQL01_S25 | D_TRACK_LISTINGS | -       |


1 rows returned in 0.01 seconds [Download](#)

Workspace: US\_1350 User: US\_1350\_SQL01\_S25

10. Drop the synonym that you created in question 9.

Rows

10



Save

Run

```
drop synonym d_track_list;
```

Results

Explain

Describe

Saved SQL

History

Synonym dropped.

0.03 seconds

Workspace: US\_1350 User: US\_1350\_SQL01\_S25